



# STK-IN4300

## Statistical Learning Methods in Data Science

Riccardo De Bin

`debin@math.uio.no`

## Outline of the lecture

- Model Assessment and Selection
  - Cross-Validation
  - Bootstrap Methods
- Methods using Derived Input Directions
  - Principal Component Regression
  - Partial Least Squares
- Shrinkage Methods
  - Ridge Regression

## Cross-Validation: $k$ -fold cross-validation

The **cross-validation** aims at estimating the **expected test error**,

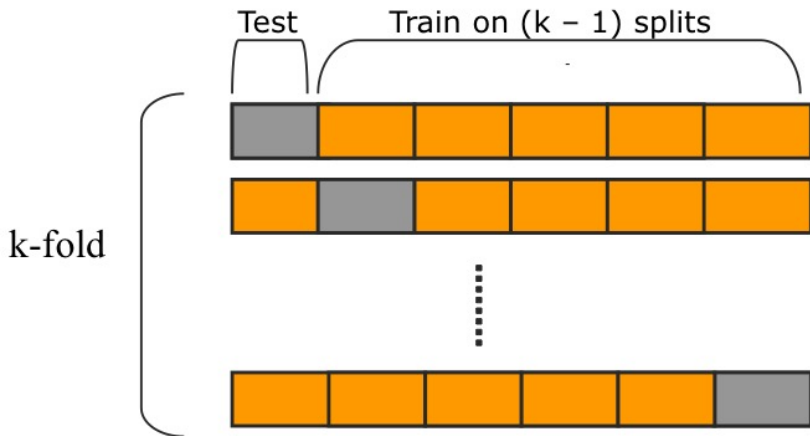
$$\text{Err} = E[L(Y, \hat{f}(X))].$$

- with **enough data**, we can **split** them in a training and test set;
- since usually it is not the case, we **mimic** this split by using the limited amount of data we have,
  - split data in  **$K$  folds**  $\mathcal{F}_1, \dots, \mathcal{F}_K$ , **approximately same size**;
  - use, in turn,  **$K - 1$  folds to train** the model (derive  $\hat{f}^{-k}(X)$ );
  - **evaluate** the model in the **remaining fold**,

$$CV(\hat{f}^{-k}) = \frac{1}{|\mathcal{F}_k|} \sum_{i \in \mathcal{F}_k} L(y_i, \hat{f}^{-k}(x_i))$$

- **estimate the expected test error** as an average,

$$CV(\hat{f}) = \frac{1}{K} \sum_{k=1}^K \frac{1}{|\mathcal{F}_k|} \sum_{i \in \mathcal{F}_k} L(y_i, \hat{f}^{-k}(x_i)) \stackrel{|\mathcal{F}_k| = \frac{N}{K}}{=} \frac{1}{N} \sum_{i=1}^N L(y_i, \hat{f}^{-k}(x_i)).$$

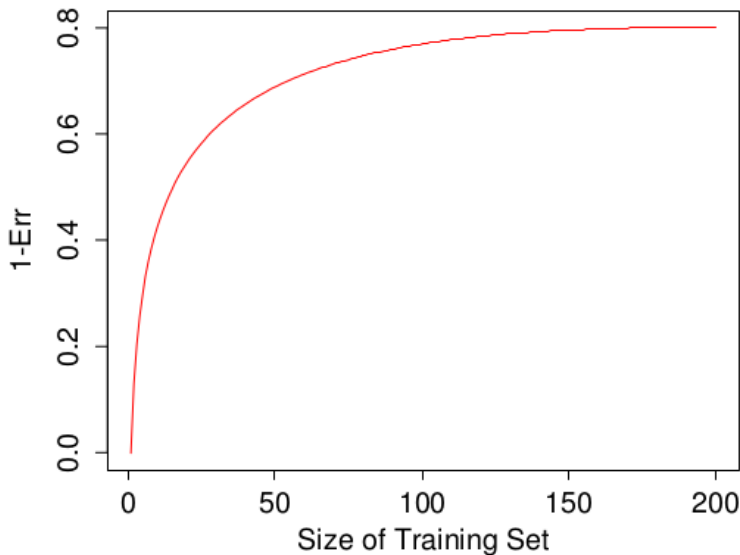
Cross-Validation:  $k$ -fold cross-validation

(figure from <http://qingkaikong.blogspot.com/2017/02/machine-learning-9-more-on-artificial.html>)

## Cross-Validation: choice of $K$

How to choose  $K$ ?

- there is no a clear solution;
- bias-variance trade-off:
  - ▶ smaller the  $K$ , smaller the variance (but larger bias);
  - ▶ larger the  $K$ , smaller the bias (but larger variance);
  - ▶ extreme cases:
    - ▶  $K = 2$ , half observations for training, half for testing;
    - ▶  $K = N$ , leave-one-out cross-validation (LOOCV);
  - ▶ LOOCV estimates the expected test error approximately unbiased;
  - ▶ LOOCV has very large variance (the “training sets” are very similar to one another);
- usual choices are  $K = 5$  and  $K = 10$ .

Cross-Validation: choice of  $K$ 

## Cross-Validation: further aspects

If we want to select a **tuning parameter** (e.g., no. of neighbours)

- train  $\hat{f}^{-k}(X, \alpha)$  for each  $\alpha$ ;
- compute  $CV(\hat{f}, \alpha) = \frac{1}{K} \sum_{k=1}^K \frac{1}{|\mathcal{F}_k|} \sum_{i \in \mathcal{F}_k} L(y_i, \hat{f}^{-k}(x_i, \alpha))$ ;
- obtain  $\hat{\alpha} = \operatorname{argmin}_{\alpha} CV(\hat{f}, \alpha)$ .

The **generalized cross-validation** (GCV),

$$GCV(\hat{f}) = \frac{1}{N} \sum_{i=1}^N \left[ \frac{y_i - \hat{f}(x_i)}{1 - \operatorname{trace}(S)/N} \right]^2$$

- is a **convenient approximation** of LOOCV for linear fitting under square loss;
- has **computational advantages**.

## Cross-Validation: the wrong and the right way to do cross-validation

Consider the following procedure:

1. find a subset of good (= most correlated with the outcome) predictors;
2. use the selected predictors to build a classifier;
3. use cross-validation to compute the prediction error.

Practical example (see R file):

- generated  $X$ , an  $[N = 50] \times [p = 5000]$  data matrix;
- generate independently  $y_i$ ,  $i = 1, \dots, 50$ ,  $y_i \in \{0, 1\}$ ;
- the true error test is 0.50;
- implementing the procedure above. What does it happen?



## Cross-Validation: the wrong and the right way to do cross-validation

Why it is **not correct**?

- Training and test sets are **NOT independent!**
- observations on the test sets are **used twice**.

**Correct** way to proceed:

- divide the sample in  $K$  folds;
- both perform variable selection and build the classifier **using observations from  $K - 1$  folds**;
  - possible choice of the tuning parameter included;
- compute the prediction error on the **remaining fold**.

## Bootstrap Methods: bootstrap

IDEA: generate pseudo-samples from the empirical distribution function computed on the original sample;

- by **sampling with replacement** from the original dataset;
- mimic new experiments.

Suppose  $Z = \{\underbrace{(x_1, y_1)}_{z_1}, \dots, \underbrace{(y_N, x_N)}_{z_N}\}$  be the **training set**:

- by sampling with replacement,  $Z_1^* = \{\underbrace{(y_1^*, x_1^*)}_{z_1^*}, \dots, \underbrace{(y_N^*, x_N^*)}_{z_N^*}\};$
- ... ..
- by sampling with replacement,  $Z_B^* = \{\underbrace{(y_1^*, x_1^*)}_{z_1^*}, \dots, \underbrace{(y_N^*, x_N^*)}_{z_N^*}\};$
- use the  $B$  bootstrap samples  $Z_1^*, \dots, Z_B^*$  to estimate **any aspect of the distribution** of a map  $S(Z)$ .

## Bootstrap Methods: bootstrap

For example, to estimate the variance of  $S(Z)$ ,

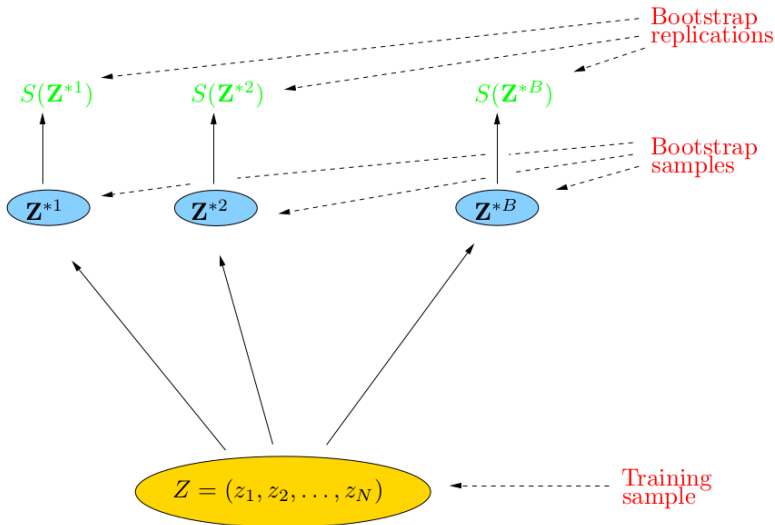
$$\widehat{\text{Var}}[S(Z)] = \frac{1}{B-1} \sum_{b=1}^B (S(Z_b^*) - \bar{S}^*)^2$$

where  $\bar{S}^* = \frac{1}{B} \sum_{b=1}^B S(Z_b^*)$ .

Note that:

- $\widehat{\text{Var}}[S(Z)]$  is the Monte Carlo estimate of  $\text{Var}[S(Z)]$  under sampling from the empirical distribution  $\hat{F}$ .

## Bootstrap Methods: bootstrap



## Bootstrap Methods: estimate prediction error

Very simple:

- generate  $B$  bootstrap samples  $Z_1^*, \dots, Z_B^*$ ;
- apply the prediction rule to each bootstrap sample to derive the predictions  $\hat{f}_b^*(x_i)$ ,  $b = 1, \dots, B$ ;
- compute the error for each point, and take the average,

$$\widehat{\text{Err}}_{\text{boot}} = \frac{1}{B} \sum_{b=1}^B \frac{1}{N} \sum_{i=1}^N L(y_i, \hat{f}_b^*(x_i)).$$

Is it correct? **NO!!!**

Again, training and test set are **NOT** independent!

## Bootstrap Methods: example

Consider a classification problem:

- two classes with the same number of observations;
- predictors and class label independent  $\Rightarrow \text{Err} = 0.5$ .

Using the 1-nearest neighbour:

- if  $y_i \in Z_b^* \rightarrow \hat{\text{Err}} = 0$ ;
- if  $y_i \notin Z_b^* \rightarrow \hat{\text{Err}} = 0.5$ ;

Therefore,

$$\widehat{\text{Err}}_{\text{boot}} = 0 \times \Pr[Y_i \in Z_b^*] + 0.5 \times \underbrace{\Pr[Y_i \notin Z_b^*]}_{0.368} = 0.184$$

## Bootstrap Methods: why 0.368

$\Pr[\text{observation } i \text{ does not belong to the bootstrap sample } b] = 0.368$

Since

$$\Pr[Z_{b[j]}^* \neq y_i] = \frac{N-1}{N},$$

is true for each position  $[j]$ , then

$$\Pr[Y_i \notin Z_b^*] = \left(\frac{N-1}{N}\right)^N \xrightarrow{N \rightarrow \infty} e^{-1} \approx 0.368,$$

Consequently,

$\Pr[\text{observation } i \text{ is in the bootstrap sample } b] \approx 0.632.$

## Bootstrap Methods: correct estimate prediction error

Note:

- each bootstrap sample has  $N$  observations;
- some of the original observations are included more than once;
- some of them (in average,  $0.368N$ ) are not included at all;
  - these are not used to compute the predictions;
  - they can be used as a test set,

$$\widehat{\text{Err}}^{(1)} = \frac{1}{N} \sum_{i=1}^N \frac{1}{|C_{[-i]}|} \sum_{b \in C_{[-i]}} L(y_i, \hat{f}_b^*(x_i))$$

where  $C_{[-i]}$  is the set of indices of the bootstrap samples which do not contain the observation  $i$  and  $|C_{[-i]}|$  denotes its cardinality.



## Bootstrap Methods: 0.632 bootstrap

Issue:

- the **average number of unique observations** in the bootstrap sample is  $0.632N \rightarrow$  not so far from  $0.5N$  of 2-fold CV;
- similar **bias** issues of 2-fold CV;
- $\widehat{\text{Err}}^{(1)}$  slightly **overestimates** the prediction error.

To solve this, the **0.632 bootstrap** estimator has been developed,

$$\widehat{\text{Err}}^{(0.632)} = 0.368 \bar{\text{err}} + 0.632 \widehat{\text{Err}}^{(1)}$$

- in practice it works well;
- in case of **strong overfit**, it can **break down**;
  - ▶ consider again the previous classification problem example;
  - ▶ with 1-nearest neighbour,  $\bar{\text{err}} = 0$ ;
  - ▶  $\widehat{\text{Err}}^{(0.632)} = 0.632 \widehat{\text{Err}}^{(1)} = 0.632 \times 0.5 = 0.316 \neq 0.5$ .

## Bootstrap Methods: 0.632+ bootstrap

Further improvement, 0.632+ bootstrap:

- based on the no-information error rate  $\gamma$ ;
- $\gamma$  takes into account the amount of overfitting;
- $\gamma$  is the error rate if predictors and response were independent;
- computed by considering all combinations of  $x_i$  and  $y_i$ ,

$$\hat{\gamma} = \frac{1}{N} \sum_{i=1}^N \frac{1}{N} \sum_{i'=1}^N L(y_i, \hat{f}(x_{i'})).$$

## Bootstrap Methods: 0.632+ bootstrap

The quantity  $\hat{\gamma}$  is used to estimate the **relative overfitting rate**,

$$\hat{R} = \frac{\widehat{\text{Err}}^{(1)} - \overline{\text{err}}}{\hat{\gamma} - \overline{\text{err}}},$$

which is then use in the **0.632+ bootstrap estimator**,

$$\widehat{\text{Err}}^{(0.632+)} = (1 - \hat{w}) \overline{\text{err}} + \hat{w} \widehat{\text{Err}}^{(1)},$$

where

$$\hat{w} = \frac{0.632}{1 - 0.368 \hat{R}}.$$

## Methods using Derived Input Directions: summary

- Principal Components Regression
- Partial Least Squares

## Principal Component Regression: singular value decomposition

Consider the **singular value decomposition** (SVD) of the  $N \times p$  (standardized) input matrix  $X$ ,

$$X = UDV^T$$

where:

- $U$  is the  $N \times p$  **orthogonal** matrix whose columns span the **column space** of  $X$ ;
- $D$  is a  $p \times p$  **diagonal** matrix, whose diagonal entries  $d_1 \geq d_2 \geq \dots \geq d_p \geq 0$  are the **singular values** of  $X$ ;
- $V$  is the  $p \times p$  **orthogonal** matrix whose columns span the **row space** of  $X$ .

## Principal Component Regression: principal components

Simple algebra leads to

$$X^T X = V D^2 V^T,$$

the **eigen decomposition** of  $X^T X$  (and, up to a constant  $N$ , of the sample covariance matrix  $S = X^T X / N$ ).

Using the eigenvectors  $v_j$  (columns of  $V$ ), we can define the **principal components** of  $X$ ,

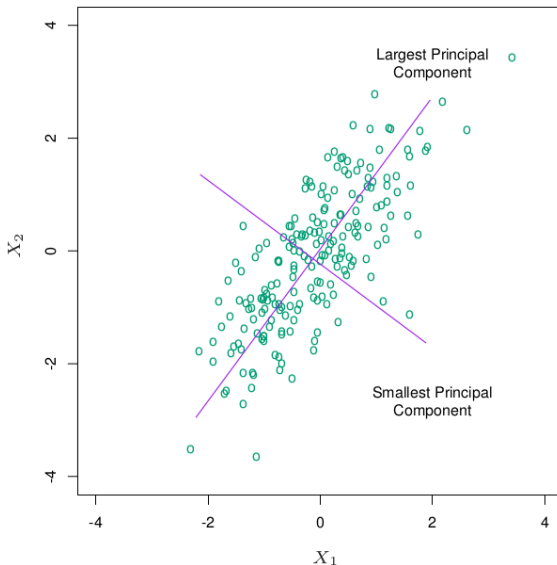
$$z_j = X v_j.$$

- the **first principal component**  $z_1$  has the **largest** sample **variance** (among all linear combinations of the columns of  $X$ );

$$\text{Var}(z_1) = \text{Var}(X v_1) = \frac{d_1^2}{N}$$

- since  $d_1 \geq \dots \geq d_p \geq 0$ , then  $\text{Var}(z_1) \geq \dots \geq \text{Var}(z_p)$ .

## Principal Component Regression: principal components



## Principal Component Regression: principal components

Principal component regression (PCR):

- use  $M \leq p$  principal components as **input**;
- regress  $y$  on  $z_1, \dots, z_M$ ;
- since the **principal components** are **orthogonal**,

$$\hat{y}_{\text{pcr}}(M) = \bar{y} + \sum_{m=1}^M \hat{\theta}_m z_m,$$

where  $\hat{\theta}_m = \langle z_m, y \rangle / \langle z_m, z_m \rangle$ .

Since  $z_m$  are linear combinations of  $x_j$ ,

$$\hat{\beta}_{\text{pcr}}(M) = \sum_{m=1}^M \hat{\theta}_m v_m.$$



## Principal Component Regression: remarks

Note that:

- PCR can be used in **high-dimensions**, as long as  $M < n$ ;
- idea: **remove** the directions with **less** information;
- if  $M = N$ ,  $\hat{\beta}_{\text{pcr}}(M) = \hat{\beta}_{\text{OLS}}$ ;
- $M$  is a **tuning parameter**, may be chosen via cross-validation;
- **shrinkage effect** (clearer later);
- principal component are scale dependent, it is **important to standardize  $X$ !**

## Partial Least Squares: idea

Partial least square (PLS) is based on an idea **similar** to PCR:

- construct a set of linear combinations of  $X$ ;
- **PCR** only uses  $X$ , **ignoring**  $y$ ;
- in PLS we want to **also** consider the information on  $y$ ;
- as for PCR, it is important to **first standardize**  $X$ .

## Partial Least Squares: algorithm

1. standardize each  $x_j$ , set  $\hat{y}^{[0]} = \bar{y}$  and  $x_j^{[0]} = x_j$ ;
2. For  $m = 1, 2, \dots, p$ ,
  - (a)  $z_m = \sum_{j=1}^p \hat{\varphi}_{mj} x_j^{[m-1]}$ , with  $\hat{\varphi}_{mj} = \langle x_j^{[m-1]}, y \rangle$ ;
  - (b)  $\hat{\theta}_m = \langle z_m, y \rangle / \langle z_m, z_m \rangle$ ;
  - (c)  $\hat{y}^{[m]} = \hat{y}^{[m-1]} + \hat{\theta}_m z_m$ ;
  - (d) orthogonalize each  $x_j^{[m-1]}$  with respect to  $z_m$ ,

$$x_j^{[m]} = x_j^{[m-1]} - \left( \frac{\langle z_m, x_j^{[m-1]} \rangle}{\langle z_m, z_m \rangle} \right) z_m, \quad j = 1, \dots, p;$$

3. output the sequence of fitted vectors  $\{\hat{y}^{[m]}\}_1^p$ .

## Partial Least Squares: step by step

First step:

- (a) compute the **first PLS direction**,  $z_1 = \sum_{j=1}^p \hat{\varphi}_{1j} x_j$ ,
  - based on the **relation between each  $x_j$  and  $y$** ,  $\hat{\varphi}_1 = \langle x_j, y \rangle$ ;
- (b) estimate the related **regression coefficient**,  $\hat{\theta}_1 = \frac{\langle z_1, y \rangle}{\langle z_1, z_1 \rangle} = \frac{\bar{z}_1 \bar{y}}{z_1^2}$ ;
- (c) model after the first iteration:  $\hat{y}^{[1]} = \bar{y} + \hat{\theta}_1 z_1$ ;
- (d) **orthogonalize**  $x_1, \dots, x_p$  w.r.t.  $z_1$ ,  $x_j^{[2]} = x_j - \left( \frac{\langle z_1, x_j \rangle}{\langle z_1, z_1 \rangle} \right) z_1$ ;

We are now ready for the second step ...

## Partial Least Squares: step by step

... using  $x_j^{[2]}$  instead of  $x_j$ :

- (a) compute the **second** PLS direction,  $z_2 = \sum_{j=1}^p \hat{\varphi}_{2j} x_j^{[2]}$ ,
  - based on the relation between each  $x_j^{[2]}$  and  $y$ ,  $\hat{\varphi}_2 = \langle x_j^{[2]}, y \rangle$ ;
- (b) estimate the related regression coefficient,  $\hat{\theta}_2 = \frac{\langle z_2, y \rangle}{\langle z_2, z_2 \rangle}$ ;
- (c) model after the second iteration:  $\hat{y}^{[2]} = \bar{y} + \hat{\theta}_1 z_1 + \hat{\theta}_2 z_2$ ;
- (d) orthogonalize  $x_1^{[2]}, \dots, x_p^{[2]}$  w.r.t.  $z_2$ ,
$$x_j^{[2]} = x_j^{[2]} - \left( \frac{\langle z_2, x_j^{[2]} \rangle}{\langle z_2, z_2 \rangle} \right) z_2;$$

and so on, **until the  $M \leq p$  step**  $\rightarrow M$  derived inputs.

## Partial Least Squares: PLS versus PCR

Differences:

**PCR** the **derived input directions** are the **principal components** of  $X$ , constructed by looking at the variability of  $X$ ;

**PLS** the input directions take into consideration **both** the **variability** of  $X$  and the **correlation** between  $X$  and  $y$ .

Mathematically:

**PCR**  $\max_{\alpha} \text{Var}(X\alpha)$ , s.t.

- $\|\alpha\| = 1$  and  $\alpha^T S v_{\ell} = 0$ ,  $\ell = 1, \dots, M - 1$ ;

**PLS**  $\max_{\alpha} \text{Cor}^2(y, X\alpha) \text{Var}(X\alpha)$ , s.t.

- $\|\alpha\| = 1$  and  $\alpha^T S \varphi_{\ell} = 0$ ,  $\forall \ell < M$ .

In practice, the **variance tends to dominate** → similar results!

## Ridge Regression: historical notes

When two predictors are **strongly correlated**  $\rightarrow$  **collinearity**;

- in the extreme case of **linear dependency**  $\rightarrow$  **super-collinearity**;
- in the case of super-collinearity,  $X^T X$  is **not invertible** (not full rank);

Hoerl & Kennard (1970):  $X^T X \rightarrow X^T X + \lambda I_p$ , where  $\lambda > 0$  and

$$I_p = \begin{pmatrix} 1 & 0 & \dots & 0 \\ 0 & 1 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 1 \end{pmatrix}.$$

With  $\lambda > 0$ ,  $(X^T X + \lambda I_p)^{-1}$  exists.

## Ridge Regression: estimator

Substituting  $X^T X$  with  $X^T X + \lambda I_p$  in the LS estimator,

$$\hat{\beta}_{\text{ridge}}(\lambda) = (X^T X + \lambda I_p)^{-1} X^T y.$$

Alternatively, the ridge estimator can be seen as the minimizer of

$$\sum_{i=1}^N (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2,$$

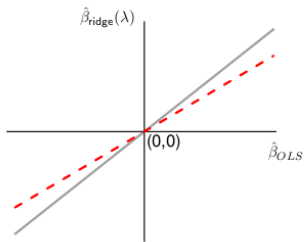
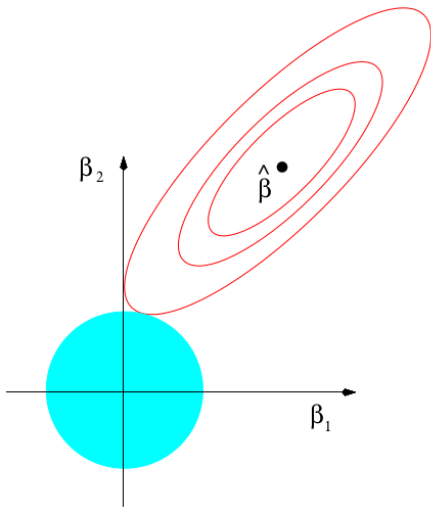
subject to  $\sum_{j=1}^p \beta_j^2 \leq t$ .

Which is the same as

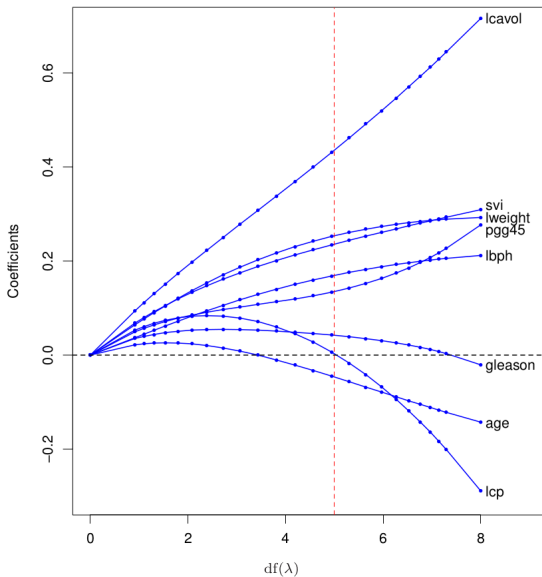
$$\hat{\beta}_{\text{ridge}}(\lambda) = \operatorname{argmin}_{\beta} \left\{ \sum_{i=1}^N (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p \beta_j^2 \right\}.$$



## Ridge Regression: visually



## Ridge Regression: visually



## Ridge Regression: remarks

Note:

- ridge solution is **not equivariant under scaling**  $\rightarrow X$  must be **standardized** before applying the minimizer;
- the intercept is **not involved** in the penalization;
- Bayesian interpretation:
  - $Y_i \sim N(\beta_0 + x_i^T \beta, \sigma^2)$ ;
  - $\beta \sim N(0, \tau^2)$ ;
  - $\lambda = \sigma^2 / \tau^2$ ;
  - $\hat{\beta}_{\text{ridge}}(\lambda)$  as the posterior mean.

## Ridge Regression: bias

$$\begin{aligned} E[\hat{\beta}_{\text{ridge}}(\lambda)] &= E[(X^T X + \lambda I_p)^{-1} X^T y] \\ &= E[(I_p + \lambda(X^T X)^{-1})^{-1} \underbrace{(X^T X)^{-1} X^T y}_{\hat{\beta}_{\text{LS}}}] \\ &= \underbrace{(I_p + \lambda(X^T X)^{-1})^{-1}}_{w_\lambda} E[\hat{\beta}_{\text{LS}}] \\ &= w_\lambda \beta \implies E[\hat{\beta}_{\text{ridge}}(\lambda)] \neq \beta \text{ for } \lambda > 0. \end{aligned}$$

- $\lambda \rightarrow 0, E[\hat{\beta}_{\text{ridge}}(\lambda)] \rightarrow \beta$ ;
- $\lambda \rightarrow \infty, E[\hat{\beta}_{\text{ridge}}(\lambda)] \rightarrow 0$  (without intercept);
- due to correlation,  $\lambda_a > \lambda_b \not\Rightarrow |\hat{\beta}_{\text{ridge}}(\lambda)| > |\hat{\beta}_{\text{ridge}}(\lambda)|$ .

## Ridge Regression: variance

Consider the variance of the ridge estimator,

$$\begin{aligned}\text{Var}[\hat{\beta}_{\text{ridge}}(\lambda)] &= \text{Var}[w_{\lambda} \hat{\beta}_{\text{LS}}] \\ &= w_{\lambda} \text{Var}[\hat{\beta}_{\text{LS}}] w_{\lambda}^T \\ &= \sigma^2 w_{\lambda} (X^T X)^{-1} w_{\lambda}^T.\end{aligned}$$

Then,

$$\begin{aligned}\text{Var}[\hat{\beta}_{\text{LS}}] - \text{Var}[\hat{\beta}_{\text{ridge}}(\lambda)] &= \sigma^2 [(X^T X)^{-1} - w_{\lambda} (X^T X)^{-1} w_{\lambda}^T] \\ &= \sigma^2 w_{\lambda} [(I_p + \lambda (X^T X)^{-1}) (X^T X)^{-1} (I_p + \lambda (X^T X)^{-1})^T - (X^T X)^{-1}] w_{\lambda}^T \\ &= \sigma^2 w_{\lambda} [(X^T X)^{-1} + 2\lambda (X^T X)^{-2} + \lambda^2 (X^T X)^{-3} - (X^T X)^{-1}] w_{\lambda}^T \\ &= \sigma^2 w_{\lambda} [2\lambda (X^T X)^{-2} + \lambda^2 (X^T X)^{-3}] w_{\lambda}^T > 0\end{aligned}$$

(since all terms are quadratic and therefore positive)

$$\implies \text{Var}[\hat{\beta}_{\text{ridge}}(\lambda)] \leq \text{Var}[\hat{\beta}_{\text{LS}}]$$

## Ridge Regression: degrees of freedom

Note that the ridge solution is a **linear combination** of  $y$ , as the least squares one:

- $\hat{y}_{\text{LS}} = \underbrace{X(X^T X)^{-1} X^T}_H y \longrightarrow df = \text{trace}(H) = p;$
- $\hat{y}_{\text{ridge}} = \underbrace{X(X^T X + \lambda I_p)^{-1} X^T}_{H_\lambda} y \longrightarrow df(\lambda) = \text{trace}(H_\lambda);$ 
  - $\text{trace}(H_\lambda) = \sum_{i=1}^p \frac{d_j^2}{d_j^2 + \lambda};$
  - $d_j$  is the diagonal element of  $D$  in the SVD of  $X$ ;
  - $\lambda \rightarrow 0, df(\lambda) \rightarrow p;$
  - $\lambda \rightarrow \infty, df(\lambda) \rightarrow 0.$

## Ridge Regression: more about shrinkage

Recall the SVD decomposition  $X = UDV^T$ , and the properties

$$U^T U = I_p = V^T V.$$

$$\begin{aligned}\hat{\beta}_{\text{LS}} &= (X^T X)^{-1} X^T y \\ &= (VDU^T UDV^T)^{-1} VDU^T y \\ &= (VD^2 V^T)^{-1} VDU^T y \\ &= VD^{-2} V^T VDU^T y \\ &= VD^{-2} DU^T y\end{aligned}$$

$$\begin{aligned}\hat{y}_{\text{LS}} &= X\hat{\beta}_{\text{LS}} \\ &= UDV^T VD^{-2} DU^T y \\ &= UDD^{-2} DU^T y \\ &= UU^T y\end{aligned}$$

## Ridge Regression: more about shrinkage

$$\begin{aligned}\hat{\beta}_{\text{ridge}} &= (X^T X + \lambda I_p)^{-1} X^T y \\ &= (VDU^T UDV^T + \lambda I_p)^{-1} VDU^T y \\ &= (VD^2 V^T + \lambda VV^T)^{-1} VDU^T y \\ &= V(D^2 + \lambda I_p)^{-1} V^T VDU^T y \\ &= V(D^2 + \lambda I_p)^{-1} U^T y\end{aligned}$$

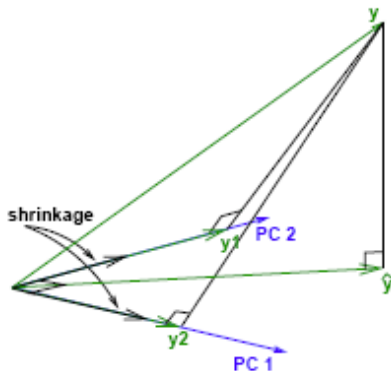
$$\begin{aligned}\hat{y}_{\text{ridge}} &= X \hat{\beta}_{\text{ridge}} \\ &= UDV^T V(D^2 + \lambda I_p)^{-1} U^T y \\ &= UV^T V D^2 (D^2 + \lambda I_p)^{-1} U^T y \\ &= U \underbrace{D^2 (D^2 + \lambda I_p)^{-1}}_{\sum_{j=1}^p \frac{d_j^2}{d_j^2 + \lambda}} U^T y\end{aligned}$$

So:

- **small** singular values  $d_j$  correspond to directions of the column space of  $X$  with **low** variance;
- ridge regression **penalizes the most** these directions.



## Ridge Regression: more about shrinkage



$N$ -dimensional sample space  
based on principle components

$$\beta_1 = \|y1\| / \|pc1\| \quad \beta_1^{ridge} = \frac{\|pc1\|^2 \beta_1}{\|pc1\|^2 + \lambda}$$

$$\beta_2 = \|y2\| / \|pc2\| \quad \beta_2^{ridge} = \frac{\|pc2\|^2 \beta_2}{\|pc2\|^2 + \lambda}$$

(picture from <https://onlinecourses.science.psu.edu/stat857/node/155/>)

## References |

HOERL, A. E. & KENNARD, R. W. (1970). Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics* **12**, 55–67.