

STK-IN4300

Statistical Learning Methods in Data Science

Riccardo De Bin

debin@math.uio.no

Gradient Boosting: from the last lecture

In the last lecture:

- boosting as implementation of “wisdom of the crowds”;
- repeatedly apply a weak learner to modification of the data;
- from AdaBoost to gradient boosting;
- forward stagewise additive modelling;
- importance of shrinkage.

Outline of the lecture

- Gradient Boosting
 - review
 - L_2 boosting with linear learner
- Likelihood-based Boosting
 - introduction
- Tree-based boosting

Gradient Boosting: general gradient boosting

The general algorithm for the gradient boosting:

1. initialize the estimate, e.g., $f_0(x) = 0$;
2. for $m = 1, \dots, m_{\text{stop}}$,
 - 2.1 compute the negative gradient vector,
$$u_m = - \left. \frac{\partial L(y, f(x))}{\partial f(x)} \right|_{f(x) = \hat{f}_{m-1}(x)};$$
 - 2.2 fit the base learner to the negative gradient vector, $h_m(u_m, x)$;
 - 2.3 update the estimate, $f_m(x) = f_{m-1}(x) + \nu h_m(u_m, x)$.
3. final estimate, $\hat{f}_{m_{\text{stop}}}(x) = \sum_{m=1}^{m_{\text{stop}}} \nu h_m(u_m, x)$

Gradient Boosting: L_2 boosting with linear learner

The L_2 Boost algorithm (linear learner, L_2 loss):

1. initialize the estimate, e.g., $\hat{\beta}_0 = 0$;
2. for $m = 1, \dots, m_{\text{stop}}$,
 - 2.1 compute the **negative gradient vector**,
$$u_m = - \left. \frac{\partial L(y, f(X, \beta))}{\partial f(X, \beta)} \right|_{\beta = \hat{\beta}_{m-1}};$$
 - 2.2 fit the **base learner** to the negative gradient vector,
$$b_m(u_m, X) = \nu(X^T X)^{-1} X^T u_m;$$
 - 2.3 update the estimate, $\hat{\beta}_m = \hat{\beta}_{m-1} + b_m(u_m, x)$.
3. final estimate, $\hat{f}_{m_{\text{stop}}}(x) = X^T \hat{\beta}_m$.

 L_2 boosting with linear learner: **properties**

Consider a **linear** learner \mathcal{S} (e.g., least square). Then

Proposition 1 (Bühlmann & Yu, 2003): The **eigenvalues** of the L_2 Boost operator \mathcal{B}_m are

$$\{1 - (1 - \lambda_k)^{m_{\text{stop}}+1}, k = 1, \dots, N\}.$$

If $\mathcal{S} = \mathcal{S}^T$ (i.e., symmetric), then \mathcal{B}_m can be diagonalized with an orthonormal transformation,

$$\mathcal{B}_m = U D_m U^T, \quad D_m = \text{diag}(1 - (1 - \lambda_k)^{m_{\text{stop}}+1})$$

where $U U^T = U^T U = I$.

 L_2 boosting with linear learner: **boosting operator**

Consider the regression model $y_i = f(x_i) + \epsilon_i$, $i = 1, \dots, N$,

- ϵ_i i.i.d. with $E[\epsilon_i] = 0$, $\text{Var}[\epsilon_i] = \sigma^2$.
- linear learner $\mathcal{S} : \mathbb{R}^N \rightarrow \mathbb{R}^N$ ($\mathcal{S}y = \hat{y}$);

Note that:

- $\hat{f}_m(x) = \hat{f}_{m-1}(x) + \mathcal{S}u_m$;
- $u_m = y - \hat{f}_{m-1}(x) = u_{m-1} - \mathcal{S}u_{m-1} = (I - \mathcal{S})u_{m-1}$;
- iterating, $u_m = (I - \mathcal{S})^m$, $m = 1, \dots, m_{\text{stop}}$.

Because $\hat{f}_m(x) = \mathcal{S}y$, then $\hat{f}_{m_{\text{stop}}}(x) = \sum_{m=0}^{m_{\text{stop}}} \mathcal{S}(I - \mathcal{S})^m y$, i.e.,

$$\hat{f}_{m_{\text{stop}}}(x) = \underbrace{(I - (I - \mathcal{S})^{m+1})}_{\text{boosting operator } \mathcal{B}_m} y.$$

 L_2 boosting with linear learner: **properties**

We can now compute:

- $\text{bias}^2(m, \mathcal{S}; f) = N^{-1} \sum_{i=1}^N (E[\hat{f}_m(x_i)] - f)^2$
$$= N^{-1} f^T U \text{diag}((1 - \lambda_k)^{2m+2}) U^T f;$$
- $\text{Var}(m, \mathcal{S}; \sigma^2) = N^{-1} \sum_{i=1}^N (\text{Var}[\hat{f}_m(x_i)])$
$$= \sigma^2 N^{-1} \sum_{i=1}^N (1 - (1 - \lambda_k)^{m+1})^2;$$

and

- $\text{MSE}(m, \mathcal{S}; f, \sigma^2) = \text{bias}^2(m, \mathcal{S}; f) + \text{Var}(m, \mathcal{S}; \sigma^2)$.


L₂ boosting with linear learner: properties

Assuming $0 < \lambda_k \leq 1$, $k = 1, \dots, N$, note that:

- $\text{bias}^2(m, \mathcal{S}; f)$ decays exponentially fast for m increasing;
- $\text{Var}(m, \mathcal{S}; \sigma^2)$ increases exponentially slow for m increasing;
- $\lim_{m \rightarrow \infty} \text{MSE}(m, \mathcal{S}; f, \sigma^2) = \sigma^2$;
- if $\exists k : \lambda_k < 1$ (i.e., $\mathcal{S} \neq I$), then $\exists m : \text{MSE}(m, \mathcal{S}; f, \sigma^2) < \sigma^2$;
- if $\forall k : \lambda_k < 1$, $\frac{\mu_k}{\sigma^2} > \frac{1}{(1-\lambda_k)^2} - 1$, then $\text{MSE}_{\mathcal{B}_m} < \text{MSE}_{\mathcal{S}}$, where $\mu = U^T f$ (μ represents f in the coordinate system of the eigenvectors of \mathcal{S}).

(for the proof, see Bühlmann & Yu, 2003, Theorem 1)

L₂ boosting with linear learner: properties



L₂ boosting with linear learner: properties

About $\frac{\mu_k}{\sigma^2} > \frac{1}{(1-\lambda_k)^2} - 1$:

- a large left side means that f is relatively complex compared with the noise level σ^2 ;
- a small right side means that λ_k is small, i.e. the learner shrinks strongly in the direction of the k -th eigenvector;
- therefore, to have boosting bringing improvements:
 - there must be a large signal to noise ratio;
 - the value of λ_k must be sufficiently small;



use a weak learner!!!

L₂ boosting with linear learner: properties

There is a further interesting theorem in Bühlmann & Yu (2003),

Theorem 2: Under the assumption seen till here and $0 < \lambda_k \leq 1$, $k = 1, \dots, N$, and assuming that $E[|\epsilon_1|^p] < \infty$ for $p \in \mathbb{N}$,

$$N^{-1} \sum_{i=1}^N E[(\hat{f}_m(x_i) - f(x_i))^p] = E[\epsilon_1^p] + O(e^{-Cm}), \quad m \rightarrow \infty$$

where $C > 0$ does not depend on m (but on N and p).

This theorem can be used to argue that boosting for classification is resistant to overfitting (for $m \rightarrow \infty$, exponentially small overfitting).

Gradient Boosting: boosting in high-dimensions

The boosting algorithm is working in **high-dimension frameworks**:


- forward stagewise additive modelling;
- at each step, **only one dimension** (component) of X is updated at each iteration;
- in a **parametric** setting, only one $\hat{\beta}_j$ is updated;
- an additional step in which it is decided **which component to update** must be computed at each iteration.

Gradient Boosting: component-wise L_2 Boost with linear learner

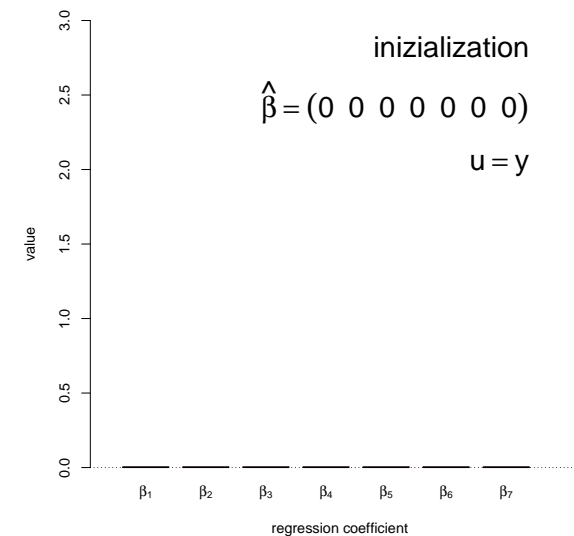
Component-wise L_2 Boost algorithm:

1. initialize the estimate, e.g., $\hat{\beta} = (0, \dots, 0)$;
2. for $m = 1, \dots, m_{\text{stop}}$,
 - compute the negative gradient vector, $u = - \left. \frac{\partial L(y, f(x, \beta))}{\partial f(x, \beta)} \right|_{\beta = \hat{\beta}}$ for the **j -th component only**;
 - fit the base learner to the negative gradient vector, $\hat{h}(u, x_j)$;
 - **select the best update j^*** (usually that minimizes the loss);
 - include the shrinkage factor, $\hat{b}_j = \nu \hat{h}(u, x_j)$;
 - update the estimate, $\hat{\beta}_{j*} = \hat{\beta}_{j*} + \hat{b}_{j*}$.
3. final estimate, $\hat{f}_{m_{\text{stop}}}(x) = X^T \hat{\beta}^{[m_{\text{stop}}]}$ (for linear regression).

Boosting: minimization of the loss function




Boosting: parameter estimation






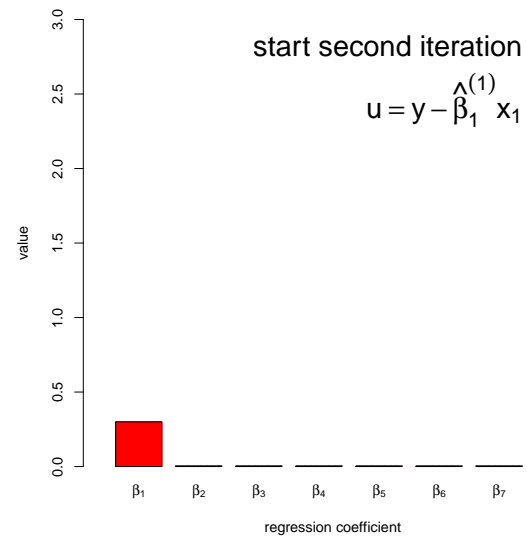
Boosting: parameter estimation




Boosting: parameter estimation



Boosting: parameter estimation




Boosting: parameter estimation







Boosting: parameter estimation



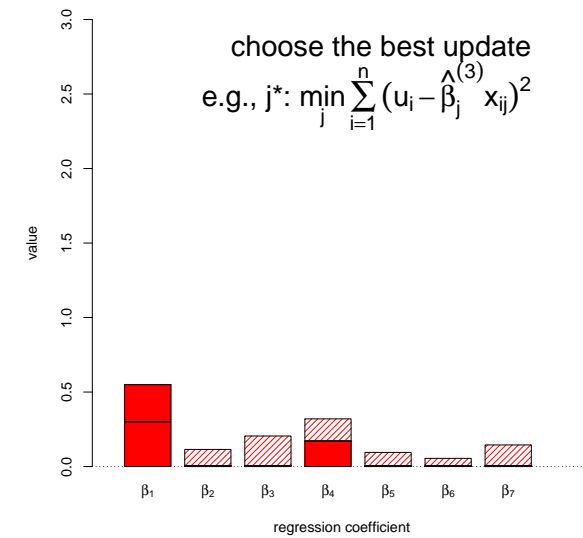
Boosting: parameter estimation



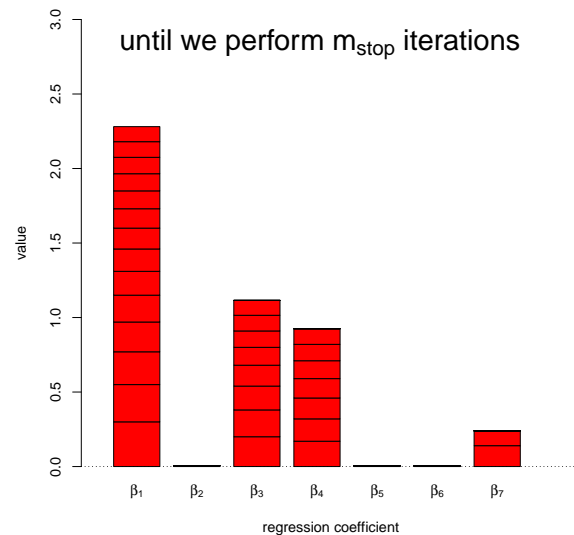
Boosting: parameter estimation



Boosting: parameter estimation



Boosting: parameter estimation



Boosting: tuning parameters

- The update step is regulated by the **shrinkage parameter** ν ;
- as long as its **magnitude is reasonable**, the choice of the penalty parameter does **not** influence the procedure;
- the choice of the **number of iterations** m_{stop} is highly relevant;
- m_{stop} (complexity parameter) influences **variable selection properties** and model sparsity;
- m_{stop} controls the amount of **shrinkage**;
 - m_{stop} **too small** results in a model which is **not able** to describe the data variability;
 - an **excessively large** m_{stop} causes **overfitting** and causes the selection of irrelevant variables.
- there is no standard approach \rightarrow **repeated cross-validation** (Seibold et al., 2016).

Likelihood-based Boosting: introduction

A different version of boosting is the so-called **likelihood-based boosting** (Tutz & Binder, 2006):

- based on the concept of **GAM** as well;
- loss function as a **negative log-likelihood**;
- uses standard statistical tools (**Fisher scoring**, basically a Newton-Raphson algorithm) to minimize the loss function;
- likelihood-based boosting and gradient boosting are **equal only** in **Gaussian** regression (De Bin, 2016).

Likelihood-based Boosting: algorithm

The simplest implementation of the likelihood-based boosting is **BoostR**, based on the ridge estimator:

Algorithm: BoostR

Step 1: Initialization. $\hat{\beta}_{(0)} = (X^T X + \lambda I_p)^{-1} X^T y$, $\hat{\mu}_{(0)} = X \hat{\beta}_{(0)}$.

Step 2: Iteration. For $m = 1, 2, \dots$ apply ridge regression to the model for residuals

$$y - \hat{\mu}_{(m-1)} = X \beta^R + \varepsilon,$$


$$\text{yielding solutions } \hat{\beta}_{(m)}^R = (X^T X + \lambda I_p)^{-1} X^T (y - \hat{\mu}_{(m-1)}), \hat{\mu}_{(m)}^R = X \hat{\beta}_{(m)}^R.$$

The new estimate is obtained by $\hat{\mu}_{(m)} = \hat{\mu}_{(m-1)} + \hat{\mu}_{(m)}^R$.

see also Tutz & Binder (2007).

In the rest of the lecture we will give the **general idea** and see its implementation as a special case of gradient boosting.

Likelihood-based Boosting: introduction




Following the statistical interpretation of boosting:

maximize the **log-likelihood** $\ell(\beta)$ (equivalently, $-\ell(\beta)$ is the **loss function** to minimize);

prediction \rightarrow **shrinkage**
aim at $\hat{\beta}_{shrink}$, not $\hat{\beta}_{MLE}$;

best solution is “between” 0 and $\hat{\beta}_{MLE}$.

Likelihood-based Boosting: introduction



starting point...
maximize a log-likelihood...



Newton-Raphson method
(or Fisher scoring).

Basic idea:

- apply Newton-Raphson;
- stop early enough to end in $\hat{\beta}_{shrink}$ and not in $\hat{\beta}_{MLE}$.

Likelihood-based Boosting: Newton-Raphson

General Newton-Raphson step:

$$\hat{\beta}^{[m]} = \hat{\beta}^{[m-1]} + \left(-\ell_{\beta\beta}(\beta) \Big|_{\beta=\hat{\beta}^{[m-1]}} \right)^{-1} \ell_{\beta}(\beta) \Big|_{\beta=\hat{\beta}^{[m-1]}},$$

where:

- $\ell_{\beta}(\beta) = \frac{\partial \ell(\beta)}{\partial \beta}$;
- $\ell_{\beta\beta}(\beta) = \frac{\partial^2 \ell(\beta)}{\partial \beta^T \partial \beta}$.

For convenience, let us rewrite the general step as

$$\underbrace{\hat{\beta}^{[m]} - \hat{\beta}^{[m-1]}}_{\text{improvement at step } m} = 0 + \left(-\ell_{\beta\beta}(\beta) \Big|_{\beta=\hat{\beta}^{[m-1]}} \right)^{-1} \ell_{\beta}(\beta) \Big|_{\beta=\hat{\beta}^{[m-1]}}.$$

Likelihood-based Boosting: Newton-Raphson

Control the Newton-Raphson algorithm:

- we need to force the estimates to be **between 0 and $\hat{\beta}_{MLE}$** ;
 - we need to **be able to stop at $\hat{\beta}_{shrink}$** .
- \Rightarrow we need smaller “controlled” improvements.


Solution: **penalize the log-likelihood!**

- $p\ell(\beta) \leftarrow \ell(\beta) - \frac{1}{2} \lambda \|\beta\|_2^2$;
- $p\ell_{\beta}(\beta) \leftarrow \ell_{\beta}(\beta) - \lambda \|\beta\|_1$;
- $p\ell_{\beta\beta}(\beta) \leftarrow \ell_{\beta\beta}(\beta) - \lambda$;

Now the general step is:

$$\underbrace{\hat{\beta}^{[m]} - \hat{\beta}^{[m-1]}}_{\text{improvement at step } m} = \left(-\ell_{\beta\beta}(\beta) \Big|_{\beta=\hat{\beta}^{[m-1]}} + \lambda \right)^{-1} \ell_{\beta}(\beta) \Big|_{\beta=\hat{\beta}^{[m-1]}}$$

Likelihood-based Boosting: visualization



As long as λ is 'big enough', the boosting learning path is going to hit $\hat{\beta}_{shrink}$.

We **must stop** at that point: the number of boosting iterations (m_{stop}) is crucial!

Likelihood-based Boosting: likelihood-based vs gradient

In the **likelihood-based boosting** we:

- repeatedly implement the first step of **Newton-Raphson**;
- update at each step **estimates** and likelihood.

Small improvements:

- parabolic approximation**;
- fit the **negative gradient** on the data by a base-learner (e.g., least-square estimator)

$$\hat{\beta}^{[m]} - \hat{\beta}^{[m-1]} = (X^T X + \lambda)^{-1} X^T \underbrace{\frac{\partial \ell(\eta(\beta, X))}{\partial \eta(\beta, X)}}_{\text{negative gradient}} \bigg|_{\hat{\beta}^{[m-1]}}$$

Likelihood-based Boosting: likelihood-based vs gradient

Substituting

$$\nu = (X^T X + \lambda)^{-1} X^T X$$

one obtains the expression of the **L2Boost** for (generalized) linear models seen before,

$$\hat{\beta}^{[m]} - \hat{\beta}^{[m-1]} = \nu (X^T X)^{-1} X^T \frac{\partial \ell(\eta(\beta, X))}{\partial \eta(\beta, X)} \bigg|_{\hat{\beta}^{[m-1]}}$$

- gradient boosting** is a **much more general** algorithm;
- likelihood-based boosting and gradient boosting are **equal in Gaussian** regression because the log-likelihood is a parabola;
- both have a **componentwise version**.

Likelihood-based Boosting: likelihood-based vs gradient

Alternatively (more correctly) we can see the likelihood-based boosting as a special case of the gradient boosting (De Bin, 2016):

- initialize $\hat{\beta} = (0, \dots, 0)$;
- for $m = 1, \dots, m_{stop}$
 - compute the negative gradient vector, $u = \frac{\partial \ell(f(x, \beta))}{\partial f(x, \beta)} \bigg|_{\beta = \hat{\beta}}$
 - compute the update,
$$\hat{b}^{LB} = \left(\frac{\partial f(x, \beta)}{\partial \beta} \bigg|_{\beta=0}^\top u \right) / \left(- \frac{\partial^2 \ell(f(x, \beta))}{\partial \beta^2} \bigg|_{\beta=0} u + \lambda \right);$$
 - update the estimate, $\hat{\beta}^{[m]} = \hat{\beta}^{[m-1]} + \hat{b}^{LB}$.
- compute the final prediction, e.g., for lin. regr. $\hat{y} = X^T \hat{\beta}^{[m_{stop}]}$

Tree-based boosting: introduction

The **base (weak) learner** in a boosting algorithm can be a **tree**:

- largely used in practice;
- very **powerful and fast** algorithm;
- R package **XGBoost**;
- we **lose part** of the statistical interpretation.

Tree-based boosting: algorithm

Algorithm 10.3 Gradient Tree Boosting Algorithm.

1. Initialize $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$.

2. For $m = 1$ to M :

(a) For $i = 1, 2, \dots, N$ compute

$$r_{im} = - \left[\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}.$$

(b) Fit a regression tree to the targets r_{im} giving terminal regions R_{jm} , $j = 1, 2, \dots, J_m$.


(c) For $j = 1, 2, \dots, J_m$ compute

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma).$$

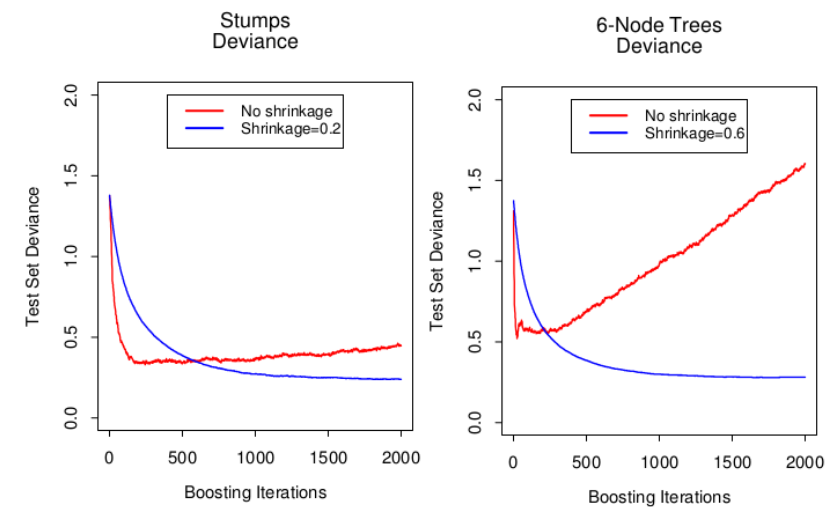
(d) Update $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$.

3. Output $\hat{f}(x) = f_M(x)$.


Tree-based boosting: importance of “weakness”



Tree-based boosting: importance of “shrinkage”



Tree-based boosting: comparison



References I

- BÜHLMANN, P. & YU, B. (2003). Boosting with the L_2 loss: regression and classification. *Journal of the American Statistical Association* **98**, 324–339.
- DE BIN, R. (2016). Boosting in Cox regression: a comparison between the likelihood-based and the model-based approaches with focus on the R-packages CoxBoost and mboost. *Computational Statistics* **31**, 513–531.
- SEIBOLD, H., BERNAU, C., BOULESTEIX, A.-L. & DE BIN, R. (2016). On the choice and influence of the number of boosting steps. Tech. Rep. 188, University of Munich.
- TUTZ, G. & BINDER, H. (2006). Generalized additive modeling with implicit variable selection by likelihood-based boosting. *Biometrics* **62**, 961–971.
- TUTZ, G. & BINDER, H. (2007). Boosting ridge regression. *Computational Statistics & Data Analysis* **51**, 6044–6059.