

STK-IN4300

Statistical Learning Methods in Data Science

Riccardo De Bin

debin@math.uio.no

Ensemble Learning: introduction

With **ensemble learning** we denote methods which:

- apply **base learners** to the data;
- **combine** the results of these learner.

Examples:

- bagging;
- random forests;
- boosting,
 - in boosting the learners **evolves** over time.

Outline of the lecture

- Ensemble Learning
 - Introduction
 - Boosting and regularization path
 - The “bet on sparsity” principle
- High-Dimensional Problems: $p \gg N$
 - When p is much larger than N
 - Computational short-cuts when $p \gg N$
 - Supervised Principal Component

Ensemble Learning: boosting and regularization path

Consider the following algorithm,

Algorithm 16.1 *Forward Stagewise Linear Regression.*

1. Initialize $\check{\alpha}_k = 0$, $k = 1, \dots, K$. Set $\varepsilon > 0$ to some small constant, and M large.
 2. For $m = 1$ to M :
 - (a) $(\beta^*, k^*) = \arg \min_{\beta, k} \sum_{i=1}^N \left(y_i - \sum_{l=1}^K \check{\alpha}_l T_l(x_i) - \beta T_k(x_i) \right)^2$.
 - (b) $\check{\alpha}_{k^*} \leftarrow \check{\alpha}_{k^*} + \varepsilon \cdot \text{sign}(\beta^*)$.
 3. Output $f_M(x) = \sum_{k=1}^K \check{\alpha}_k T_k(x)$.
-


(note its similarities with the boosting algorithm)

Ensemble Learning: boosting and regularization path

In comparison with lasso:

- initialization ($\check{\alpha}_k = 0, k = 1, \dots, K$) $\longleftrightarrow \lambda = \infty$;
- for small values of M :
 - some $\check{\alpha}_k$ are not updated \longleftrightarrow coefficients “forced” to be 0;
 - $\check{\alpha}_k^{[0]} \leq \check{\alpha}_k^{[M]} \leq \check{\alpha}_k^{[\infty]} \longleftrightarrow$ shrinkage;
 - M inversely related to λ ;
- for M large enough ($M = \infty$ in boosting) and $K < N$,
 $\check{\alpha}_k^{[M]} = \hat{\alpha}_{LS} \longleftrightarrow \lambda = 0$

Ensemble Learning: boosting and regularization path



Ensemble Learning: boosting and regularization path

If

- all the basis learners T_K are mutually uncorrelated,

then

- for $\varepsilon \rightarrow 0$ and $M \rightarrow \infty$, such that $\varepsilon M \rightarrow t$,

Algorithm 16.1 gives the lasso solutions with $t = \sum_k |\alpha_k|$.

In general component-wise boosting and lasso do not provide the same solution:

- in practice, often similar in terms of prediction;
- for $\varepsilon (\nu) \rightarrow 0$ boosting (and forward stagewise in general) tends to the path of the least angle regression algorithm;
- lasso can also be seen as a special case of least angle.

Ensemble Learning: boosting and regularization path


Consider 1000 Gaussian distributed variables:

- strongly correlated ($\rho = 0.95$) in blocks of 20;
- uncorrelated blocks;
- one variable with effect on the outcome for each block;
- effects generated from a standard Gaussian.


Moreover:

- added Gaussian noise;
- noise-to-signal ratio = 0.72.

Ensemble Learning: boosting and regularization path



Ensemble Learning: boosting and regularization path



Ensemble Learning: the “bet on sparsity” principle

We consider:

- L_1 -type of penalty (shrinkage, variable selection);
- L_2 -type of penalty (shrinkage, computationally easy);
- boosting's stagewise forward strategy minimizes something close to a L_1 penalized loss function;
- step-by-step minimization.

Can we characterize situations where one is preferable to the other?

Ensemble Learning: the “bet on sparsity” principle

Consider the following framework:

- 50 observations;
- 300 independent Gaussian variables.


Three scenarios:

- all 300 variables are relevant;
- only 10 out of 300 variables are relevant;
- 30 out of 300 variables are relevant.

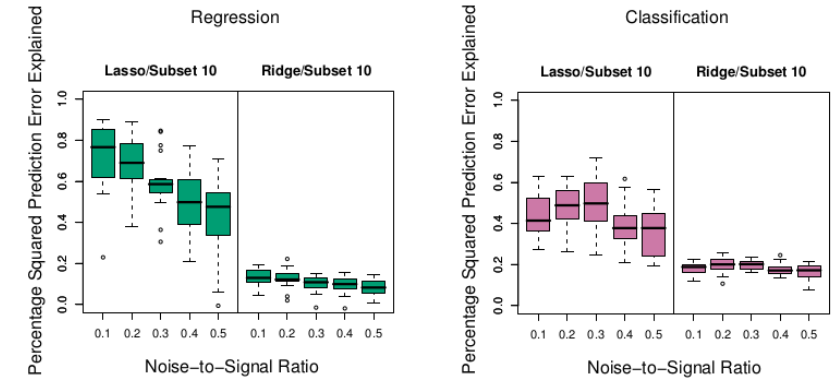
Outcome:

- regression (added standard Gaussian noise);
- classification (from an inverse-logit transformation of the linear predictor)

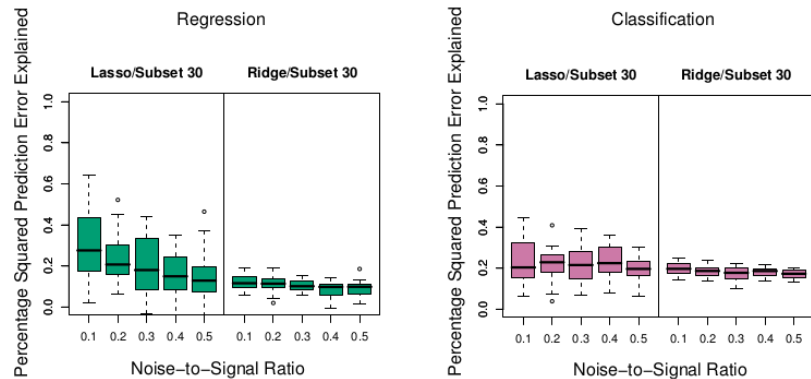
Ensemble Learning: the “bet on sparsity” principle



Ensemble Learning: the “bet on sparsity” principle



Ensemble Learning: the “bet on sparsity” principle



Ensemble Learning: the “bet on sparsity” principle

This means that:

- lasso performs **better** than ridge in **sparse** contexts;
- ridge gives better results if there are **several** relevant variables with **small effects**;
- **anyway**, in the dense case, the model **does not** explain a lot,
 - not enough data to estimate correctly several coefficients;

↓
“bet on sparsity”
=

“use a procedure that does well in sparse problems, since no procedure does well in dense problems”

Ensemble Learning: the “bet on sparsity” principle

The **degree of sparseness** depends on:

- the **unknown mechanism** generating the data;
 - it depends on the number of relevant variables;
- **size of the training set**;
 - larger sizes allow estimating denser models;
- **noise-to-signal ratio**;
 - smaller NTR → denser models (same as before);
- **size of the dictionary**;
 - more base learners, potentially sparser models.

High-Dimensional Problems: when p is much larger than N

The case $p \gg N$ (number of variable much larger than the number of observations):

- **very important** in current applications,
 - e.g., in a common genetic study, $p \approx 23000$, $N \approx 100$;
- concerns about **high variance** and **overfitting**;
- **highly regularized approaches** are common:
 - lasso;
 - ridge;
 - boosting;
 - elastic-net;
 - ...

High-Dimensional Problems: when p is much larger than N

Consider the following example:

- $N = 100$;
- p variables from standard Gaussians with $\rho = 0.2$,
 - (i) $p = 10$ (ii) $p = 100$ (iii) $p = 1000$.
- response from

$$Y = \sum_{j=1}^p X_j \beta_j + \sigma \epsilon;$$

- signal to noise ratio $\text{Var}[E(Y|X)]/\sigma^2 = 2$;
- true β from a standard Gaussian;

High-Dimensional Problems: when p is much larger than N


As a consequence, averaging over 100 replications:

- p_0 , the number of significant β (i.e., $\beta : |\hat{\beta}/\hat{\text{se}}| > 2$)
 - (i) $p_0 = 9$ (ii) $p_0 = 33$ (iii) $p_0 = 331$.

Consider 3 values of the penalty λ :

- $\lambda = 0.001$, which corresponds to
 - (i) d.o.f. = 20 (ii) d.o.f. = 99 (iii) d.o.f. = 99;
- $\lambda = 100$, which corresponds to
 - (i) d.o.f. = 9 (ii) d.o.f. = 35 (iii) d.o.f. = 87;
- $\lambda = 1000$, which corresponds to
 - (i) d.o.f. = 2 (ii) d.o.f. = 7 (iii) d.o.f. = 43.

High-Dimensional Problems: when p is much larger than N



High-Dimensional Problems: when p is much larger than N

Remarks:

- with $p = 20$, ridge regression can find the relevant variables;
 - the covariance matrix can be estimated;
- moderate shrinkage works better in the middle case, in which we can find some non-zero effects;
- with $p = 1000$, there is no hope to find the relevant variables, and it is better to shrink down everything;
 - no possibility to estimate the covariance matrix.

High-Dimensional Problems: computational short-cuts when $p \gg N$

Consider the single-value decomposition of X ,

$$\begin{aligned} X &= UDV^T \\ &= RV^T \end{aligned}$$

Where

- V is a $p \times N$ matrix with orthonormal columns;
- U is a $N \times N$ orthonormal matrix;
- D is a diagonal matrix with elements $d_1 \geq d_2 \geq \dots \geq d_N \geq 0$;
- R is a $N \times N$ with rows r_i .

High-Dimensional Problems: computational short-cuts when $p \gg N$

Theorem (Hastie et al., 2009, page 660):

Let $f^*(r_i) = \theta_0 + r_i^T \theta$ and consider the optimization problems:

$$\begin{aligned} (\hat{\beta}_0, \hat{\beta}) &= \operatorname{argmin}_{\beta_0, \beta \in \mathbb{R}^p} \sum_{i=1}^N L(y_i, \beta_0 + x_i \beta) + \lambda \beta^T \beta; \\ (\hat{\theta}_0, \hat{\theta}) &= \operatorname{argmin}_{\theta_0, \theta \in \mathbb{R}^N} \sum_{i=1}^N L(y_i, \theta_0 + r_i \theta) + \lambda \theta^T \theta. \end{aligned}$$

Then $\hat{\beta}_0 = \hat{\theta}_0$ and $\hat{\beta} = V \hat{\theta}$.

High-Dimensional Problems: computational short-cuts when $p \gg N$

Note:

- we can replace the p -dim vectors x_i with the N -dim vectors r_i ;
- same penalization, but with way fewer predictors;
- the N -dimensional solution $\hat{\theta}$ is transformed back to the p -dimensional $\hat{\beta}$ by simple matrix multiplication;
- it only works for linear models;
- it only works for quadratic penalties.
- a $O(p^3)$ problem is reduced to a $O(p^2N)$ problem;
 - relevant for $p > N$.

High-Dimensional Problems: computational short-cuts when $p \gg N$

Example (ridge regression):

Consider the estimate of a ridge regression,

$$\hat{\beta} = (X^T X + \lambda I)^{-1} X^T y.$$

Replacing X with RV^T , we obtain

$$\hat{\beta} = V(R^T R + \lambda I)^{-1} R^T y,$$

i.e.,

$$\hat{\beta} = V \hat{\theta},$$

where

$$\hat{\theta} = (R^T R + \lambda I)^{-1} R^T y.$$

High-Dimensional Problems: computational short-cuts when $p \gg N$

Note:

- it cannot be applied to lasso;
- the short-cut is particularly relevant for finding the best λ via cross-validation;
- it can be shown that one need to construct R only once,
 - use the same R for each of the CV-folds.

High-Dimensional Problems: supervised principal component

Algorithm 18.1 Supervised Principal Components.

1. Compute the standardized univariate regression coefficients for the outcome as a function of each feature separately.
 2. For each value of the threshold θ from the list $0 \leq \theta_1 < \theta_2 < \dots < \theta_K$:
 - (a) Form a reduced data matrix consisting of only those features whose univariate coefficient exceeds θ in absolute value, and compute the first m principal components of this matrix.
 - (b) Use these principal components in a regression model to predict the outcome.
 3. Pick θ (and m) by cross-validation.
-

High-Dimensional Problems: supervised principal component


Note:

- in the first step, the variables **univariately most associated** with the outcome are selected:
 - the measure depends on the **nature of the outcome**;
 - include all associations larger than a **threshold θ** ;
 - may include **highly correlated** variables.
- in step 2, **perform PCR** on the reduced variable space:
 - use the **first m components**;
 - assure **shrinkage**.
- both θ and m **must be computed** by cross-validation:
 - 2-dimension** tuning parameter.

High-Dimensional Problems: supervised principal component


Example (survival analysis with microarray data):

- data from Rosenwald et al. (2002);
- input: 7399 gene expressions of 240 patients;
- response: survival time (potentially right censored);
- divided in a training (160 patients) and (80) test set.



High-Dimensional Problems: supervised principal component

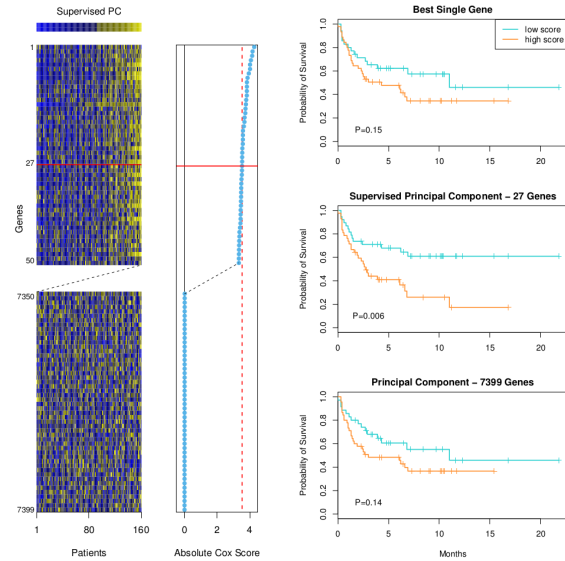
Survival Function



High-Dimensional Problems: supervised principal component




High-Dimensional Problems: supervised principal component



STK-IN4300: lecture 12

33/ 43

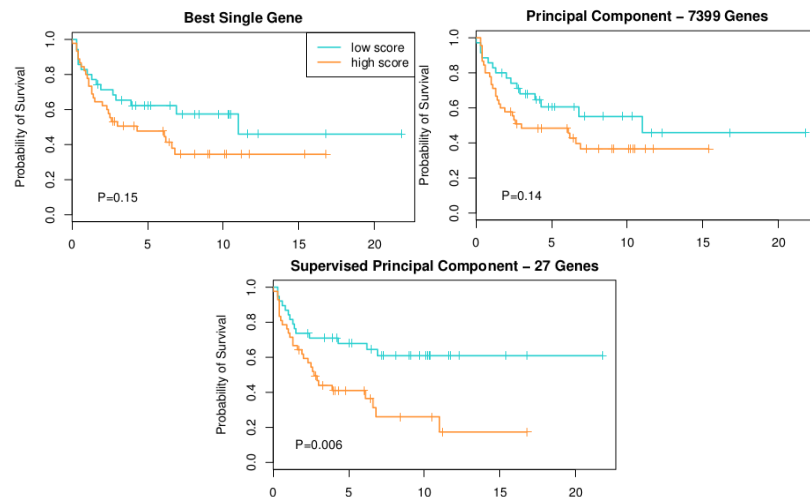
High-Dimensional Problems: supervised principal component



STK-IN4300: lecture 12

34/ 43

High-Dimensional Problems: supervised principal component



STK-IN4300: lecture 12

35/ 43

High-Dimensional Problems: supervised-PC and latent-variable modelling

Consider a latent-variable model,

$$Y = \beta_0 + \beta_1 U + \epsilon$$

where

$$X_j = \alpha_{0j} + \alpha_{1j} U + \epsilon_j \quad \text{if } j \in \mathcal{P}$$

- ϵ and ϵ_j have **mean 0** and they are **independent** of the other variables in their respective models;
- X_j are **independent** of U if $j \notin \mathcal{P}$.

STK-IN4300: lecture 12

36/ 43

High-Dimensional Problems: supervised-PC and latent-variable modelling

Supervised-PC can be seen as a method to fit this kind of model:

- step 1: **identify** the $j \in \mathcal{P}$;
 - on average β_1 is not zero only if $\alpha_{1j} \neq 0$;
- step 2a: **estimate** α_{0j} and α_{1j}
 - natural if $\epsilon_j \sim N(0; \sigma^2)$;
- step 2b: estimate β_0 and β_1 (**fit the model**).

High-Dimensional Problems: supervised-PC and partial least square

Thresholded PLS can be seen as a **noisy version** of supervised-PC,

- first **PLS variate**,
- $$z = \sum_{j \in \mathcal{P}} \langle y, x_j \rangle x_j;$$
- **supervised principal component direction**,
- $$u = \frac{1}{d^2} \sum_{j \in \mathcal{P}} \langle y, x_j \rangle x_j.$$

Set $p_1 = |\mathcal{P}|$. It can be shown (Bair & Tibshirani, 2004) that, for $N, p_1, p \rightarrow \infty$, s.t. $p_1/N \rightarrow 0$

$$z = u + O_p(1)$$

$$\hat{u} = u + O_p(\sqrt{p_1/N})$$

where u is the true latent variable.

High-Dimensional Problems: supervised-PC and partial least square

Supervised-PC versus **partial least square**:


- both aim at considering both **large variation** and **correlation with the outcome**;
- supervised-PC **removes** those which are not relevant;
- partial least square **downgrades** them



“thresholded PLS”:

- apply PLS on **only** the variables selected by supervised-PC.

High-Dimensional Problems: supervised-PC and partial least square



High-Dimensional Problems: pre-conditioning

Supervised-PC can also be used to improve lasso performance, through pre-conditioning

- compute \hat{y}_i the supervised-PC prediction for each observation;
- apply lasso using \hat{y}_i instead of y_i as outcome;
 - using all variables, not only those selected by supervised-PC.

The idea is to remove the noise, so lasso is not affected by the large number of noisy variables.

References I

- BAIR, E. & TIBSHIRANI, R. (2004). Semi-supervised methods to predict patient survival from gene expression data. *PLoS Biology* 2, e108.
- HASTIE, T., TIBSHIRANI, R. & FRIEDMAN, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference and Prediction (2nd Edition)*. Springer, New York.
- ROSENWALD, A., WRIGHT, G., CHAN, W. C., CONNORS, J. M., CAMPO, E., FISHER, R. I., GASCOYNE, R. D., MULLER-HERMELINK, H. K., SMELAND, E. B., GILTNAME, J. M. et al. (2002). The use of molecular profiling to predict survival after chemotherapy for diffuse large-b-cell lymphoma. *New England Journal of Medicine* 346, 1937–1947.

High-Dimensional Problems: pre-conditioning

